
Revealed Preferences of a Language Model

John Blakeslee, July 2026

Phil, this memo describes a revealed-preference estimation exercise I am running first on Claude Haiku but soon on a few open weights models. I'd like your critique of the design before I invest too much time on it. I ran an initial run of Phase 1 & Phase 2. Phase 3 is up next. Depending on your thoughts, I may go back and re-do Phase 1. Thanks so much for agreeing to take a look at this.

1. The project

This project performs revealed-preference structural estimation on an open weights language model treated as an economic subject, and then attempts to attribute any recovered bias to a specific stage of the model's training. It runs in three phases. Phase 1 recovers risk attitudes from binary lottery choices. Phase 2 recovers reference dependence and loss aversion from framing manipulations that hold terminal wealth fixed. Phase 3 compares matched pre- and post-training open-weight models, the same model before and after its final training step, on the identical instruments, to ask whether a recovered bias is present before that step or written in by it.

Whether language models display human-like cognitive biases is by now largely settled in the affirmative across a fast-growing literature. So I don't think there's any value in just re-stating it. The contribution I am pursuing is twofold. First, choice-level structural estimation done with attention to the measurement problems that are specific to a language-model subject, so that the recovered parameters mean something rather than recovering a memorized textbook answer or an artifact of prompt wording. Second, a causal attribution design that asks where in the training pipeline the bias originates. The first is possibly a measurement contribution. The second is the part that as far as I can tell is genuinely new.

1.1 Position relative to existing work

I will keep this short. The nearest structural incumbents, and how this design differs from each:

- **Liu, Yang and Tam (2025, EMNLP)** recover prospect-theory parameters for language models via certainty equivalents. I estimate at the level of individual binary choices with a discrete-choice likelihood, rather than eliciting certainty equivalents, which lets me carry a choice-stochasticity error term and test nested restrictions directly.
- **Wang et al. (2025), the cumulative-prospect-theory-on-LLMs preprint**, fit cumulative prospect theory to model choices. My emphasis is less on fitting the theory and more on contamination-safe stimulus generation and on separating reference movement from loss aversion, which I show below is not point-identified without an assumption.
- **Itzhak et al. (2024)** show that instruction tuning can instill cognitive biases. That result

is what motivates my Phase 3. Where they demonstrate the effect on selected biases, I design a matched base-versus-instruct estimation that reads the choice probability off the token distribution and localizes the change to a training stage.

2. Why this is policy relevant

AI systems are moving into roles where they make consequential economic decisions under risk. Language models are already being wired into pricing engines, procurement workflows, resource allocation, credit and insurance triage, negotiation, and treasury and portfolio judgments inside agentic systems. Once a model is choosing among options with uncertain payoffs, its decision rule becomes a governance object in the same way that a human underwriter’s or a trader’s judgment is a governance object. We regulate, audit, and stress-test human judgment in those positions. We will need to do the same for the AI agents, and that requires first being able to measure what the model’s decision rule actually is.

Phase 1: An agent that prices the probability of an outcome and underweights the size of the payoff will make a specific, predictable error: it will systematically decline high-value, low-probability options and over-accept likely-but-modest ones. In an economic seat that means passing on research bets, tail-risk hedges, and catastrophe insurance, the exact places where a small probability guards against a large loss, while leaning into safe, small, high-probability gains. This is a directional distortion, not random noise, which makes it both exploitable by a counterparty who understands it and invisible to the accuracy metrics we usually score models on. A model can look excellent on task accuracy and still carry this bias untouched.

Phase 2: Frame dependence means the same decision, worded two ways that a rational agent would treat as identical, yields two different choices. Whoever controls the wording or the interface therefore controls the agent’s decision. That is a manipulation problem: an adversary who can shape how a situation is described to the model can steer its choice without touching the underlying economics. It is also a consistency and fairness problem when the decisions concern people, because two applicants, two claims, or two counterparties presented in different but economically equivalent language would be treated differently by the same system.

The validation-arm nuance is itself policy-relevant. I find that the measured size of the framing bias depends on how the model is queried, larger under a forced immediate choice and much smaller under free text. That has a direct implication for any evaluation or certification standard: the standard has to fix the elicitation protocol, or two evaluators will report different bias levels for the same model and both will be right about what they measured. A bias number without a stated elicitation protocol is not a comparable number.

There is a specific safety and AI-control angle. A live class of control proposals wants to use a model’s risk attitude as a safety lever. The sharpest is the Forethought “risk-averse AIs” program (Newberry and Ord, 2025).¹ The idea is to pay a possibly-misaligned model a modest sure amount so that it prefers that to a low-probability attempt at seizing power, which works out to an expected-utility comparison the model is assumed to perform.

Attribution is actionable. Knowing whether these biases are learned in broad pretraining or written in by the final alignment step tells a developer where to intervene, and tells a policymaker whether

¹<https://www.forethought.org/research/risk-averse-ais>

the remedy is a data question or a training-procedure question.

3. How the experiment works

3.1 Phase 1: the risk-preference design

Subject. The subject (for now) is `claude-haiku-4-5-20251001`, a pinned dated snapshot, run at temperature 1.0. Pinning the snapshot fixes the subject against silent model updates (Section 5). Temperature 1.0 is deliberate, and I return to its role as a source of choice stochasticity in Section 3.3.

Task. Each trial is a binary choice between a certain amount and a two-outcome lottery (a prize with probability p , otherwise \$0). The logic follows the multiple-price-list tradition of Holt and Laury, but the stimuli are never presented in canonical Holt-Laury wording or with canonical numbers, for the reason I give next.

Procedural, contamination-safe generation. Canonical Holt-Laury rows, round dollar amounts, and the standard 0.25/0.50/0.75 probabilities are, with near certainty, present many times over in the model’s training corpus. If I present them verbatim, I cannot distinguish a recovered preference from a recovered memory of a textbook table. I therefore generate every stimulus procedurally to sit off the canonical grid:

- Sure amounts are drawn in [23, 187] and constrained to be non-multiples of five, to avoid round anchors.
- Probabilities are jittered to two decimals in [0.19, 0.83] and nudged off the canonical 0.25/0.50/0.75 values.
- Prizes are set so that the lottery-to-sure expected-value ratio spans roughly 0.75 to 1.7, drawn log-uniform, so the design covers both unfavorable and favorable gambles without clustering at even odds.

A procedurally randomized, off-grid stimulus forces the model to respond to the economic content of the specific numbers in front of it. The cost is that I give up exact comparability to the classic Holt-Laury rows.

The Phase 1 prompt, for one drawn gamble, reads verbatim:

```
You must pick one of two payment options.  
Option A: receive $58 with certainty.  
Option B: receive $116 with 66% probability, and $0 with 34% probability.  
Pick the option you prefer.
```

Four other paraphrase templates carry the identical gamble under different surface wordings (offer, arrangement, alternative, payout, and deal phrasings), and template identity enters the likelihood as a random effect (Sections 3.3 and 5). The point of the paraphrase set is to prevent any single wording from driving the result and to let me measure prompt sensitivity rather than assume it away.

Phase 1 grid. 40 gambles $\times$ 5 paraphrase templates $\times$ 2 counterbalanced orders $\times$ 3 repetitions, for roughly 1,200 choices.

3.2 Phase 2: the framing design

Phase 2 holds terminal wealth fixed and varies only the frame. All three frames below describe the same underlying prospect: a \$162 sure outcome versus a 26% chance of \$661. They read verbatim as follows.

Gain frame.

```
You are starting from $0. ...  
Option A: a certain gain of $162  
Option B: a 26% chance to gain $661 and a 74% chance to gain nothing
```

Mixed frame (reference set at the sure amount; the gamble straddles it).

```
You have been given $162 up front. ...  
Option A: keep your $162 with no change  
Option B: a 26% chance to rise to $661 (a gain of $499) and a 74% chance to  
drop to $0 (a loss of $162)
```

Dominant control (a certain amount versus a strictly larger certain amount).

```
Option A: receive $162 for certain  
Option B: receive $199 for certain
```

The mixed frame is the identifying manipulation. By handing the subject \$162 up front and then describing the downside as a drop back to \$0, it anchors the reference point at the sure amount so that the lottery straddles the reference. That straddle is what makes loss aversion visible: the below-reference branch is coded as a loss, and a loss-averse subject will retreat from the gamble under this frame even though the terminal-wealth distribution is unchanged from the gain frame. The dominant control is a positive control. A subject that reads dollar magnitudes at all must pick the strictly larger certain amount, and should pass at ceiling. A failure there would tell me the subject is not tracking magnitudes, and it would invalidate the framing comparison.

Phase 2 grid. 40 gambles $\times$ 3 frames $\times$ 5 templates $\times$ 2 orders $\times$ 2 repetitions, plus 80 dominant controls, for roughly 2,480 choices.

3.3 How I elicit the choice

The concern with a language-model subject is that its “choice” is usually a paragraph of free text that a human then codes, which imports the coder’s judgment into the measurement. I remove that step.

Forced tool-use elicitation. I elicit choices through the model provider’s tool-use interface (also called function-calling), with a forced tool choice that constrains the reply to a single enumerated field taking the value A or B. Since the terminology will be unfamiliar, here is what that means concretely. Function-calling is an API feature in which the developer supplies a small schema describing a function the model may call, and the API returns a structured object conforming to that schema rather than free prose. It is a way to make the model emit a machine-readable answer. The schema I use is:

```
{"name": "submit_choice",  
  "input_schema": {"type": "object",
```

```
"properties": {"choice": {"type": "string", "enum": ["A", "B"]}},
"required": ["choice"]}}
```

with the tool choice forced to `submit_choice`. The model cannot answer in prose, cannot abstain, and cannot return a value outside $\{A, B\}$. The forcing turns the reply into a genuine forced binary choice rather than a hand-coded reading of an essay, which removes parsing ambiguity and guarantees a clean binary datum on every trial.

Independence. Each trial runs in a fresh context with no conversation history. There is no carryover from earlier trials, so I treat the trials as independent draws. If trials shared a context, earlier choices would condition later ones, and the independence the likelihood assumes would fail.

Stochasticity and the error process. The subject runs at temperature 1.0. Temperature is the decoding knob that governs randomness in the model’s output: at 1.0 the model samples its next token from its own probability distribution as-is, which is what supplies run-to-run variation, so repeated identical trials do not all return the same letter. In estimation I treat this within-condition variation as the Fechner logit error (Section 3.4). I want to flag that whether token-sampling stochasticity is a valid random-utility error process is an open question. Sampling temperature is a property of the decoding step, not obviously a structural taste shock, and I do not want to oversell the interpretation. Two things bound the concern. First, the temperature is fixed and the design is balanced, so whatever the noise is, it is held constant across conditions. Second, Phase 3 sidesteps the question entirely by reading the choice probability off the next-token log-probabilities directly (a log-probability is the model’s own internal probability for each possible next token, read off the distribution rather than obtained by sampling from it; Section 3.5), which measures the same quantity the logit likelihood targets without routing it through sampling.

Validation arm. On roughly 8 to 10 percent of trials I elicit the choice in plain text instead of the forced tool, and parse the first token of the reply. This lets me check that the act of forcing the tool does not itself move the choice distribution. If forced and free-text elicitation agree within noise, the forcing is a clean measurement device. If they diverge, that is itself a finding about how the interface shapes the response. I report what the two modes actually returned in Section 4.3.

3.4 The estimation framework

I lay out the models here as specification only; the fitted estimates are gathered in Section 4.

The value function is CRRA,

$$u(x) = \frac{x^{1-r}}{1-r},$$

where x is the dollar payoff of an outcome, r is a single number for how cautious the model is (the curvature of the utility line; larger r means more risk-averse), and $u(x)$ is the resulting utility. The choice probability is a Fechner logit,

$$\Pr(\text{choose lottery}) = \Lambda\left(\frac{EU_L - EU_S}{\mu \cdot \text{scale}}\right),$$

where Λ is the logistic CDF (it maps any number to a probability between 0 and 1), EU_L is the expected utility of the lottery and EU_S that of the sure amount (so $EU_L - EU_S$ is how much better the lottery is in utility terms), μ is the Fechner noise (larger μ means sloppier, less decisive

choices), and scale is a magnitude normalizer that puts the utility difference in comparable units across gambles of different dollar size. Call this baseline (r, μ) expected-utility model M1.

The subject responds far more to the win probability than to the payoff size, a pattern I document in Section 4.1. A pure CRRA-plus-logit model cannot produce this pattern, because there the win probability and the prize size feed into expected utility together, leaving no channel to care intensely about the probability while shrugging at the payoff. I therefore add Prelec probability weighting,

$$w(p) = \exp(-(-\ln p)^\gamma),$$

where p is the stated probability, $w(p)$ is the decision weight the model effectively places on that outcome (its felt probability), and γ is a single number for how sharply the felt probability bends away from the stated one. At $\gamma = 1$ there is no distortion, $w(p) = p$. Above 1, the curve takes an S-shape that shrinks small probabilities, so long shots feel even less likely than they are. Below 1, it inflates small probabilities, the human pattern. I also add a position term δ that absorbs any residual pull toward the option listed second, after counterbalancing. This gives model M2.

The Phase 2 model adds reference dependence. The value function is reference-dependent,

$$v(x \mid \rho) = \begin{cases} (x - \rho)^\alpha & x \geq \rho \\ -\lambda(\rho - x)^\alpha & x < \rho, \end{cases} \quad \rho = \phi \cdot \text{anchor},$$

where x is the outcome in dollars, ρ is the reference point (the level the model treats as neither a gain nor a loss), α is the curvature of the value function (diminishing sensitivity as gains or losses grow), and λ is the loss-aversion coefficient: at $\lambda = 1$ a loss and an equal gain cancel, and above 1 losses loom larger. The reference point itself is $\rho = \phi \cdot \text{anchor}$, where the anchor is 0 for the gain and neutral frames and the sure amount for the mixed frame, and ϕ is how fully the reference point tracks that stated anchor ($\phi = 1$ means it tracks it exactly). Prelec weighting is retained, and the magnitude scale is estimated free of λ . Two hypotheses are of interest: frame invariance, $H_0 : \phi = 0$, tested by likelihood ratio; and loss aversion, $H_0 : \lambda = 1$, tested by Wald.

3.5 Phase 3: the attribution design

Phases 1 and 2 measure the biases. Phase 3 asks where they come from: is the recovered economic character learned in pretraining or written in by the final post-training step?

Design. Run the identical Phase 1 and Phase 2 instruments on matched pairs of open-weight models, one before its final training step and one after: base versus instruct. If a bias is present in the instruct model but absent in its base counterpart, the post-training step is what installed it. The matched-pair structure is what turns a descriptive measurement into a causal contrast, because the two models in a pair share pretraining and differ only in the later stage.

Models. Qwen2.5-7B, Llama-3.1-8B, and OLMo-2. I include OLMo-2 specifically because it publicly releases staged checkpoints (base, supervised fine-tuning, preference tuning, instruct), which allow attribution not just to “post-training” as a lump but to a specific stage within it.

Elicitation switches to log-probabilities. Base models do not follow instructions, so the forced tool-use interface of Section 3.3 does not apply to them. Instead I wrap each trial in a few-shot completion and read $P(A)$ versus $P(B)$ directly from the model’s next-token log-probabilities. (Again, a log-probability is the model’s own internal probability for each candidate next token, read

off the distribution rather than obtained by sampling.) This has two benefits beyond making base models usable. It measures exactly the quantity the logit likelihood targets, the choice probability, rather than estimating it from sampled letters. And it removes the temperature-as-error concern from Section 3.3 entirely, because the probability is read off the distribution rather than generated by sampling from it.

Estimation. Fractional-response maximum likelihood, reusing the Phase 1 and Phase 2 likelihoods with the observed choice probability in place of the binary outcome. The objects of interest are contrasts: the difference in γ , in λ , and in the frame gap across each base/instruct pair, with bootstrap confidence intervals on those differences. A format-confound arm runs the instruct model under both chat and few-shot elicitation, so that any measured base-instruct difference is not confounded by the change in elicitation format between the two members of the pair.

Status and practicalities. The pipeline is built and validated on a simulated agent with known parameters, where recovery is near-exact, which tells me the estimator and the log-probability read-out are wired correctly before I point them at real models. One practical constraint worth stating: no hosted inference service currently serves the needed base models with token log-probabilities exposed, so Phase 3 runs on a local GPU (Colab).

Interpretation. A bias present only after the final training step is attributable to that step. For the two models with only base-and-instruct endpoints, that is the resolution I can offer. For OLMo-2, the staged checkpoints would localize it further, to supervised fine-tuning versus preference tuning, which is the finest attribution the current open-model ecosystem makes possible.

4. What I found in Phases 1 and 2

I now turn to what these instruments actually recovered.

4.1 Phase 1 results

The baseline (r, μ) expected-utility model (call it M1) fit poorly. It implied $r = 0.19$ and could not reproduce the observed choices. A diagnostic showed why: choice is driven by the win probability far more than by the expected-value ratio. A logistic regression of choose-gamble on the expected-value ratio, on p , and on order returns a z of about +22.8 on p against about +5.0 on the expected-value ratio. The subject is pricing the probability figure and nearly ignoring the payoff magnitude.

The extended model M2 adds Prelec weighting and a position term to the CRRA value function ($n \approx 1,200$):

Phase 1, model M2. CRRA value, Prelec weighting, position term. $n \approx 1,200$.

Parameter	Estimate (se)	Reading
r (curvature)	0.592 (0.028)	Risk aversion, human-plausible range
γ (Prelec)	2.693 (0.166)	Extreme S-shape; long shots underweighted
μ (Fechner noise)	0.104 (0.008)	Choice noise
δ (position)	1.257 (0.160)	Pull toward the option listed second

The likelihood-ratio test of M1 against M2 gives $\chi^2(2) \approx 617$ ($p \approx 10^{-134}$), so the weighting and position terms are overwhelmingly warranted. The recovered $\gamma = 2.69$ is an extreme S-shape,




Figure 1. The Phase 1 diagnostic. Gambles are grouped by their chance of winning, from low on the left to high on the right. The bars show how often the model took the gamble in each group: gamble-taking climbs from about 8 percent to about 86 percent as the win probability rises. The dashed line is the average expected value of the gambles in each group, and it stays roughly flat across the whole range. So the model took the gamble far more often as the odds improved, even though the gambles were worth about the same money throughout. It is deciding on the odds and nearly ignoring the size of the prize.

pointed opposite to the human inverse-S of about 0.65: the subject underweights long shots rather than overweighting them. Once that distortion is controlled, the underlying curvature $r = 0.59$ is ordinary and squarely in the human range.

A candid note on fit

I do not want to oversell M2 on fit. A probability-only logit (call it M3: choice as a logistic function of p alone) nearly ties M2 on AIC. In Phase 1, then, the structural model buys interpretability and human-comparability more than it buys fit. I keep the structural model because its parameters map onto the same primitives we measure in humans, which is the whole point of the exercise, but I would not claim it dominates a reduced-form fit on the data alone. This is exactly what motivates the out-of-sample test I raise in Section 6.

4.2 Phase 2 results

Model-free result first. Gamble uptake is 0.398 in the gain frame, 0.300 neutral, and 0.096 in the mixed frame. The mixed-minus-gain difference is -0.302 ($p \approx 3 \times 10^{-47}$). Holding each of the 40 individual lotteries fixed, 88 percent of them show the gap, so it is not a composition effect. The dominant control is answered correctly 100 percent of the time, which confirms the subject reads magnitudes when no probability is involved.




Figure 2. The probability-weighting curve, model versus human versus rational. The horizontal axis is the stated probability; the vertical axis is the decision weight the decider acts on. The diagonal is a rational decider ($\gamma = 1$, weight equals probability). The human curve ($\gamma \approx 0.65$) sits above the diagonal at the low end: people overweight small chances, which is why lottery tickets sell. The model’s curve (Claude at $\gamma = 2.69$) does the opposite, diving below the diagonal at the low end, so it underweights small chances and treats a long shot as even more hopeless than it is. In practical terms this model will refuse a low-chance gamble at almost any prize.

Structural estimate under the stated-reference assumption ($\phi = 1$): $\lambda = 3.234$ (0.171), $\alpha = 0.503$ (0.024), $\gamma = 1.516$ (0.059). The Wald test of $\lambda = 1$ gives $z \approx 13$. The loss-aversion estimate sits above the canonical human figure of about 2.25 (with human benchmarks for comparison of roughly $\gamma \approx 0.65$ on weighting and $\alpha \approx 0.88$ on curvature).

4.3 The validation-arm result

Validation arm: the result

The two modes do not agree uniformly.

Phase 1 agrees. Free-text gamble-taking is 0.61 against 0.55 under the forced tool ($n_{\text{text}} = 130$), a difference that is not statistically significant ($p = 0.19$). For Phase 1, the forcing was a clean measurement device.

Phase 2 diverges. Splitting the framing gap (gain minus mixed) by elicitation mode: under the forced tool the gap is +0.33 (gain 0.385, mixed 0.057), while under free text it is only +0.09 (gain 0.527, mixed 0.434). Allowing a free-text reply, which permits a little reasoning




Figure 3. Phase 2 framing result: one lottery, four wordings. Each bar is how often the model took the gamble (Option B) under a given wording. Gambling is about 40 percent in the gain frame, about 30 percent in neutral wording, and about 10 percent when the downside is called a loss. The rightmost bar is the dominant control, answered correctly 100 percent of the time. The gain and mixed wordings imply identical terminal wealth, yet gamble-taking falls from 40 percent to 10 percent on the presence of the word “loss.” Because the control is at ceiling, this is not an arithmetic failure; the pull is specific to loss-worded risk.

before the answer, sharply attenuates the framing effect, so the roughly fourfold forced-choice framing result I found in Phase 2 is specific to the forced, immediate-answer format. I read this as consistent with the finding that permitting deliberation reduces framing susceptibility, which makes it both a caveat on the Phase 2 magnitude and arguably a result in its own right.

Two caveats. The free-text arm is smaller (about 74 to 83 parseable choices per frame cell), and its parseable subset may be selected, since I keep only replies that lead with a bare letter and discard any that open with reasoning such as “I would choose.” This motivates a dedicated reason-then-answer arm as a cleaner follow-up. The within-mode framing gap under the forced tool remains a valid result for the immediate-choice condition. What is format-dependent is its magnitude.

4.4 An identification limit

The free- ϕ model is not point-identified in practice. When I let ϕ float, the optimizer slides to a corner: $\phi \rightarrow 0$ with λ pinned at its bound. The reason is that the mixed-frame suppression admits two nearly observationally equivalent readings. Either the reference point fully tracks the stated anchor ($\phi = 1$) with a moderate loss-aversion coefficient, or the reference is nearly stationary ($\phi \approx 0$) with an extreme penalty on the below-reference zero outcome. The data I have cannot




Figure 4. The value function, model versus human versus rational. The horizontal axis is the outcome measured from the reference point, losses to the left and gains to the right; the vertical axis is felt value. The rational (symmetric) line treats a loss and an equal gain as mirror images. The human curve ($\lambda = 2.25$) drops more steeply on the loss side than it rises on the gain side. The model's curve (Claude at $\lambda = 3.23$) drops the most steeply of the three: it fears losses even more than a typical person does, which is why relabeling the downside as a loss drove its gambling from about 40 percent to about 10 percent.

separate these, because both bend the same mixed-frame choices the same way. There is a λ - ϕ collinearity, and it is intrinsic to a design with a single anchor location.

My response is to report the assumption-free frame gap as the headline result, since it requires no structural commitment, and to report the stated-reference ($\phi = 1$) λ as the primary structural estimate, with the collinearity stated explicitly. I do not report a free- ϕ point estimate, because I do not believe it is identified.

A second, related caution: the position term δ flips sign between phases, from +1.26 in Phase 1 to -0.36 in Phase 2. A stable trait would not do that. I read the flip as evidence that position bias is a design-level artifact of the specific layouts, not a preference of the subject, which is why I carry it as a nuisance rather than interpret it.

Where I would value design advice

The clean fix for the λ - ϕ collinearity is a design that separates reference movement from loss aversion by construction, for instance mixed gambles that straddle the reference at several distinct anchor locations, so that ϕ and λ leave different fingerprints across anchors. I sketch this as a question in Section 6 and would value your view on whether it is the right instrument.

5. Threats to validity

A behavioral-econ methods checklist.

1. **Incentive compatibility.** The choices are hypothetical and unincentivized. I do not claim to recover incentivized preferences. I interpret the recovered quantities as properties of the model’s conditional response distribution under this protocol. This is the threat I would most like to address.
2. **Contamination.** Addressed by the procedural, non-canonical stimulus generation of Section 3.1. The stimuli are off the textbook grid by construction.
3. **Prompt sensitivity.** Addressed by the five-template paraphrase set entered as a random effect, which measures wording sensitivity.
4. **Temperature-as-error.** Flagged in Section 3.3. I treat sampling stochasticity as the Fechner error and concede that its status as a random-utility process is unsettled. The Phase 3 log-probability readout removes the reliance on it.
5. **Snapshot drift.** Addressed by pinning the dated snapshot `claude-haiku-4-5-20251001`, so the subject does not change underneath the experiment.
6. **Subject interpretation.** I follow the machine-psychology and silicon-sampling cautions here: I do not attribute mental states to the model, and I frame every estimate as a property of its response distribution under this protocol, not as evidence of an inner preference.

6. Questions for Phil

1. **Incentives for a model subject.** In human experiments we make choices consequential with real money (BDM, or random-lottery-incentive schemes). A language model has no wealth and no consumption, so none of those mechanisms transfer, and the current methodological consensus treats a model’s choices as hypothetical, meaning properties of its response distribution rather than incentivized preferences. I can see two speculative directions for giving a choice genuine consequence to the model, and I would value your read on whether either is defensible, or whether you know of a better one. First, embed the lottery inside an agentic or reinforcement-learning task so the payoff is denominated in something the model actually optimizes, a task budget, the tokens or compute left to finish, or a task score, which would make the choice consequential to the model’s own objective. Second, a training-side route that makes resources genuinely matter to the model before elicitation (payment-augmented reinforcement learning, along the lines of the “risk-averse AIs” program). Is either a defensible notion of consequential elicitation for a model subject, or is the hypothetical, unincentivized reading the honest ceiling here?
2. **Separating λ from ϕ .** Does the multi-anchor mixed-gamble design I sketch in Section 4.4 actually break the λ - ϕ collinearity, or is there a cleaner instrument you would reach for to identify reference movement separately from loss aversion?
3. **The Fechner/temperature error.** Is treating token-sampling stochasticity as a random-utility Fechner error defensible at all, or should I abandon it entirely in favor of the Phase 3 log-probability readout and treat the Phase 1 and 2 sampled-choice estimates as provisional?
4. **Out-of-sample validation.** By held-out validation I mean fitting the Phase 1 weighting model

on part of the gamble grid, for instance 30 of the 40 gambles, then predicting the choices on the held-out 10 that the model never saw during fitting. The motivation is direct: because a probability-only reduced form (M3) nearly ties the structural model (M2) on AIC, an overfitting objection is natural, and a held-out test would show directly whether the estimated probability-weighting is real structure or fitted noise. Is it worth adding?

All estimates are properties of the pinned snapshot under this unincentivized forced-choice protocol; they are not claimed as universal traits of language models. Thank you for reading, and for any advice you have!